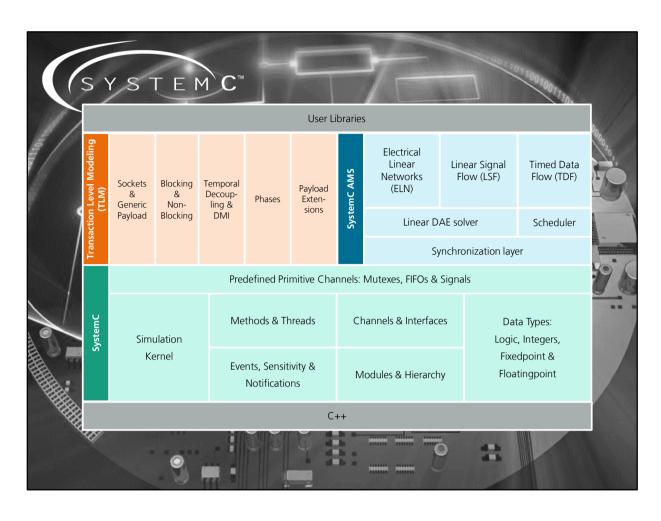


Your notes:			



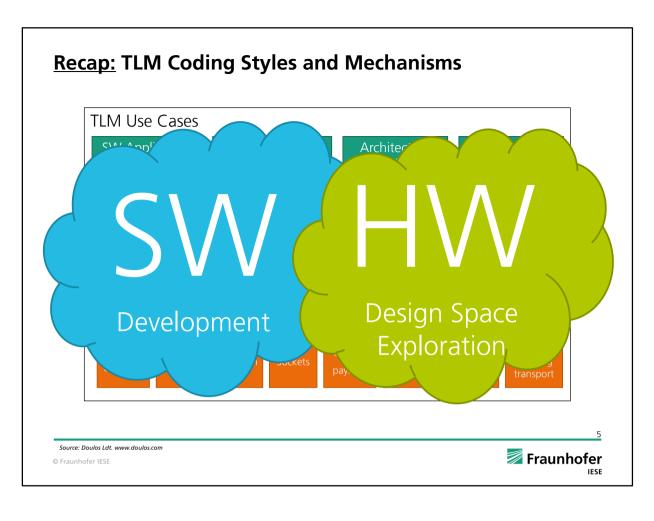
Your notes:	

			ય C™		User L	ibraries		- 1990 A	11/1
Transaction Level Modeling (TLM)	Sockets & Generic	Blocking & <u>Non-</u>	Temporal Decoup- ling &	Phases	Payload Exten-	SystemC AMS	Electrical Linear Networks (ELN)	Linear Signal Flow (LSF)	Timed Data Flow (TDF)
saction	Payload	Blocking	DMI		sions	Syste	Linear D	AE solver	Scheduler
Trans							Sy	nchronization lay	er
	Predefined Primitive Channels: Mutexes, FIFOs & Signals								
SystemC	Sim	ulation	Me	Methods & Threads			Channels & Interfaces Data Ty Logic, Int		
S S	Ke	ernel		Events, Sensitivity & Notifications		F			edpoint & atingpoint
					C	++			

Your notes:			

Recap: TLM Coding Styles and Mechanisms TLM Use Cases SW Application SW Performance Architecture Hardware Development Analysis Analysis Verification TLM 2.0 Coding Style (Just Guidelines) Loosely-Timed (LT) Multi-phase, non-blocking API Single-phase, blocking API Approximately-Timed (AT) ling Interoperability) TLM Mechanisms (Definitive API for enab Source: Doulos Ldt. www.doulos.com Fraunhofer © Fraunhofer IESE

Your notes:			



Your notes:	

Recap: Coding Styles in TLM

- Loosely-Timed (LT):
- SW
- As fast as possible
- Sufficient timing detail to boot OS and run multicore systems and to <u>develop</u> SW or drivers
- Processes can run ahead of simulation time (temporal decupling)
- Each transaction completes in one blocking function call
- Usage of Direct Memory Interface (DMI) e.g. for boot process

Approximately-Timed (AT):



- Accurate enough for performance modelling
- Sufficient for architectural HW design space exploration
- Processes run in lockstep with simulation time
- Each transaction has usually 4 timing points i.e. 4 function calls (extensible if required, also less possible); non-blocking behavior
- More detailed than LT and therefore also slower than LT

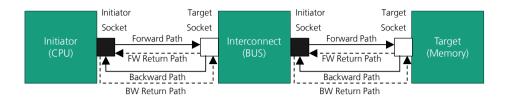
6



Your notes:	

Initiators, Targets and Interconnect





- References to the object are passed along the forward and backward paths:
 - LT uses Forward and Return Path
 - AT uses Forward, Backward, FW Return Path, BW Return Path
- AT uses non-blocking transport
- Time is handeled with Payload Event Queues (PEQs)
- Allows modelling of O-O-O Cores and Backpressure
- Base Protocol

© Fraunhofer IESE



Fraunhofer

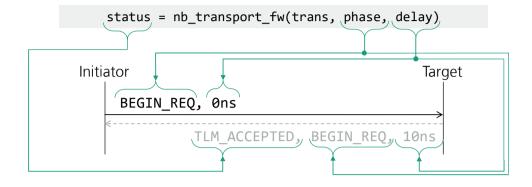
ur notes:	

		/ S T	T E I	v C™		User Li	braries	1		1011001001110	1
Transaction law Modeling	(TLM)	Sockets & Generic	Blocking & Non-	Temporal Decoup- ling &	Phases	Payload Exten-	SystemC AMS	Electrical Linear Networks (ELN)	Linear Signa Flow (LSF)	Timed Data Flow (TDF)	1000 O.
dispession	38700	Payload	Blocking	DMI		sions	Syste		AE solver	Scheduler	
		Predefined Primitive Channels: Mutexes, FIFOs & Signals									19
SystemC		Simı	ulation	Me	Methods & Threads		Channels & Interfaces			Data Types: Logic, Integers,	
		Kernel			Events, Sensitivity & Notifications		Modules & Hierarchy			Fixedpoint & Floatingpoint	
HARA						C-	++				
¥			0	. :		D. B. B. B.	h h a b a		*		

Your notes:			

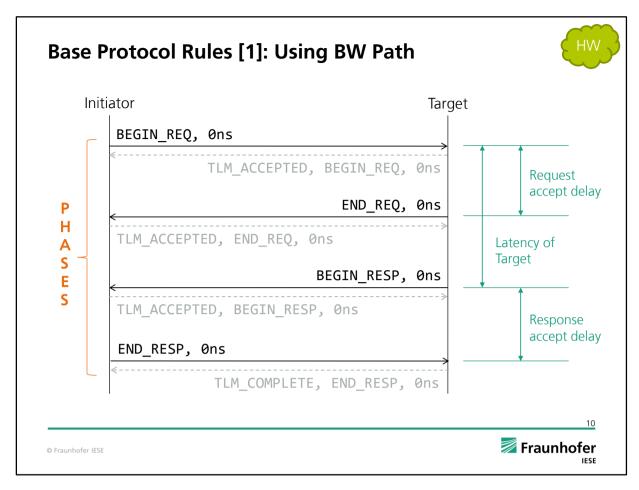
Non-Blocking Transport (AT) Base Protocol





Fraunhofer
IESE

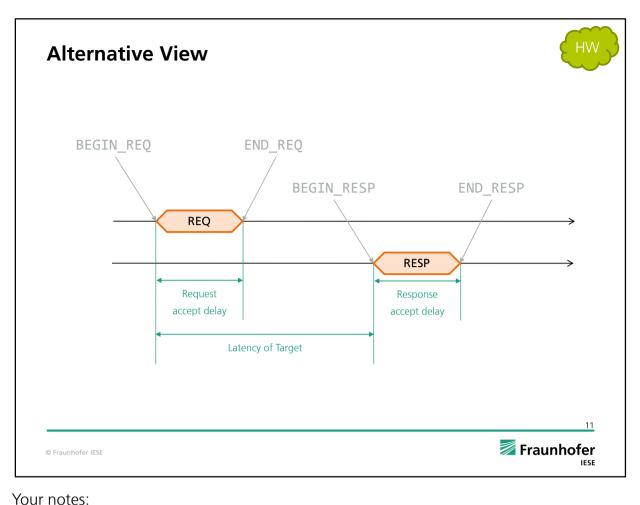
Your notes:			



An initiator, and only an initiator, may call **b_transport**. The target can return immediately, or can suspend (by calling **wait**) and return later. However, calling **wait** statements in targets should be avoided because a wait statement will result in a context switch of the simulator, which decreases the simulation speed.

The same transaction object could be reused from one call to the next, but the two calls would count as separate transactions.

Your notes:			



rour riotes.			

Base Protocol Rules: Using BW Path



- Base Protocol phases
 - BEGIN REQ \rightarrow END REQ \rightarrow BEGIN RESP \rightarrow END RESP
 - Must occur in increasing simulation time order
 - Phases must change with each call
- nb_transport_fw must not call nb_transport_bw directly and vice versa (PEQ!)
- Generic Payload memory management rules (See TLM Advanced)
- Extensions must be ignorable (See TLM Advanced)
- Target should handle mixed b_transport / nb_transport

12



ur notes:	

Base Protocol Rules: The Exclusion Rule



- Initiators and targets must honor the *Exclusion Rule*:
 - An Initiator must not send a new request (BEGIN_REQ) until it has received the END REQ from the previous transaction
 - An target must not send the next BEGIN_RESP until it has received the END RESP from the previous transactions
- The exclusion rules enable flow control like back pressure:
 - E.g. if an input buffer of a target is full the target can defer the sending of END_REQ for the transaction that filled the buffer, until the buffer has available space again
 - An RTL ready signal can be modeled by deferring END_REQ
 - However, since it is non-blocking, the target can do something else, e.g. sending

13



Your notes:			

Modelling of Backpressure



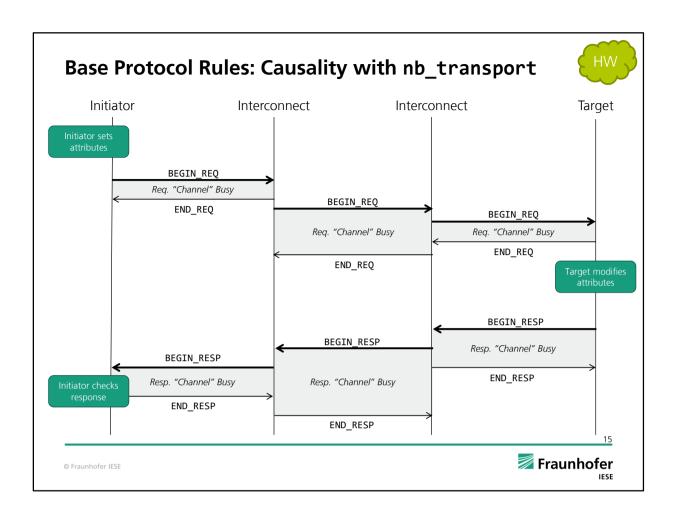


- The exclusion rules enable flow control like back pressure:
 - E.g. if an input buffer of a target is full the target can defer the sending of END_REQ for the transaction that filled the buffer, until the buffer has available space again
 - An RTL ready signal can be modeled by deferring END_REQ
 - However, since it is non-blocking, the target can do something else, e.g. sending
- Also Response exclusion rule must be honored!

14



Your notes:			



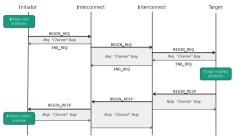
Your notes:	

Base Protocol Rules: Causality with nb_transport



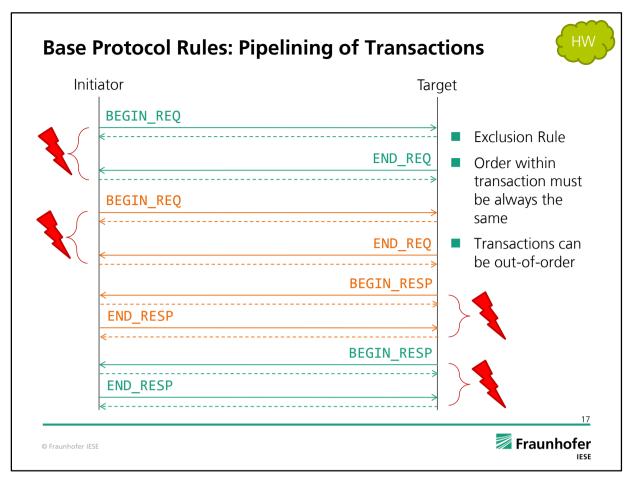
- BEGIN_REQ and BEGIN_RESP are <u>propagated</u> from end-to-end
- END_REQ and END_RESP are not propagated → they are <u>local</u> to each hop
- The END_REQ and END_RESP phases indicate that the requests and response "channels" are no longer busy and may be used for the next request or response respectively

Initiator should not check the transaction payload until it has received the response



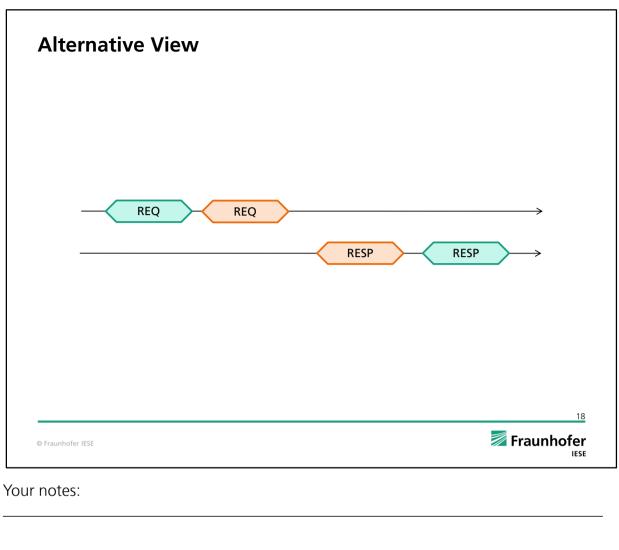


our notes:



For **nb_transport_(fw|bw)**, the timing annotation on successive calls for a given transaction must be strictly non-decreasing. For separate transactions, the mutual order is unconstrained, i.e. out-of-order transaction ordering is allowed.

Your notes:			

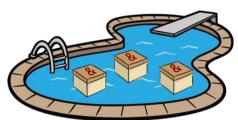


			S interest (S. Salandario)		User Li	braries	And the second s		\ \ \ \ \ \ \ \ \ \ \ \ \ \ \ \ \ \ \
Transaction Level Modeling (TLM)	Sockets & <u>Generic</u>	Blocking & Non-	Temporal Decoup- ling &	Phases	Payload Exten-	SystemC AMS	Electrical Linear Networks (ELN)	Linear Signal Flow (LSF)	Timed Data Flow (TDF)
saction	<u>Payload</u>	Blocking	DMI		sions	Syste	Linear D.	AE solver	Scheduler
Trans							Sy	nchronization l	ayer
			Pred	defined Pri	mitive Cha	nnels:	Mutexes, FIFOs &	Signals	
SystemC	Simı	ulation	Me	thods & Th	nreads	Cł	annels & Interface		Data Types: gic, Integers,
S	Kε	ernel		nts, Sensit Notificatio		М	odules & Hierarch		xedpoint & oatingpoint
					C-	++			

Your notes:	

Generic Payload Pools (a.k.a Memory Management)

- Allocating of generic payload objects consumes wall-clock time
- Idea: do not allocate for each transaction a new generic payload → Reuse
 - A *Memory Manager* handles the generic payload objects in a pool
 - Before sending a transaction a payload object is allocated from the pool
 - Modules that send (or receive) this payload object must increase a reference count (acquire), which signalizes the memory manager that this object is still in use.
 - If a module is finished with the payload object, the reference count is decreased (release)
 - If the reference count is 0 the payload is freed and goes back into the pool.
 - If all transactions in the pool are in use a new one is generated



20



Your notes:	

Memory Manager in TLM

- The memory manager is not part of the SystemC TLM Standard
- The TLM Standard provides only an interface class:

```
class tlm_mm_interface
{
    public:
    virtual void free(tlm::tlm_generic_payload*) = 0;
    virtual ~tlm_mm_interface(){}
}
```

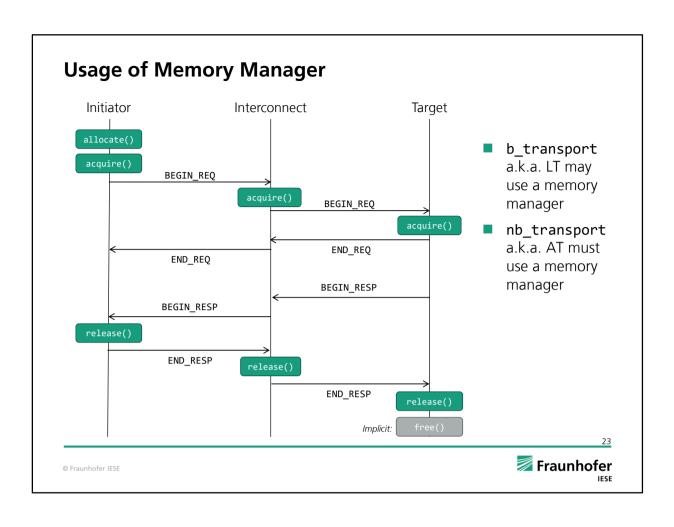
- The implementation of the memory manager is up to the end user
- Virtual Platform tools like Synopsys Platform Architect have clever implementations



Your notes:		

```
Example Memory Manager
                                                                                    Use code on github:
                                                                                   https://github.com/TUK-
                                                                        SCVP/SCVP.artifacts/tree/master/tlm_memory_manager
class MemoryManager : public tlm::tlm_mm_interface {
    private:
     unsigned int numberOfAllocations;
                                                                                         Transaction Pool
     unsigned int numberOfFrees;
                                                                                         implemented as
     std::vector<tlm::tlm_generic_payload*> freePayloads;
                                                                                         std::vector
    MemoryManager(): numberOfAllocations(∅), numberOfFrees(७) {}
    ~MemoryManager(){
         for(tlm::tlm_generic_payload* payload: freePayloads) {
                                                                                         Cleanup
             delete payload;
                                                                                         transaction pool
             numberOfFrees++;
    }
    tlm::tlm_generic_payload* MemoryManager::allocate(){
                                                                                         Allocate new
         if(freePayloads.empty()) {
                                                                                         payload
             numberOfAllocations++;
             return new tlm::tlm_generic_payload(this);
                                                                                         Reuse Payload
             tlm::tlm_generic_payload* result = freePayloads.back();
             freePayloads.pop_back();
                                                                                         and remove
             return result;
                                                                                         from pool
    }
    void free(tlm::tlm_generic_payload* payload) {
   payload->reset(); //clears all fields and extensions
                                                                                         Add transaction
                                                                                         back to the pool
         freePayloads.push_back(payload);
};
                                                                                                 Fraunhofer
© Fraunhofer IESE
```


Your notes:



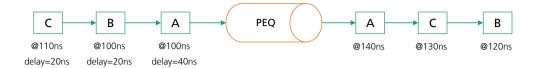
our notes:

			ST LETTER STATE OF		User Li	braries	And the second s			
Transaction Level Modeling (TLM)	Sockets & Generic	Blocking & Non-	Temporal Decoup- ling &	Phases	Payload Exten-	SystemC AMS	Electrical Linear Networks (ELN)	Linear Flow		Timed Data Flow (TDF)
saction	Payload	<u>Blocking</u>	DMI	sions	di sions	DMI sions p	AE solver		Scheduler	
Trans							Sy	nchroniz:	ation lay	er
		Predefined Primitive Channels: Mutexes, FIFOs & Signals								
SystemC	Simi	Me Simulation		Methods & Threads		Channels & Interfaces		es	Data Types: Logic, Integers,	
Ś	Kε	ernel	Events, Sensitivity & Notifications		Modules & Hierarchy		ny	Fixedpoint & Floatingpoint		
					C-	++				

Your notes:	

The Payload Event Queue



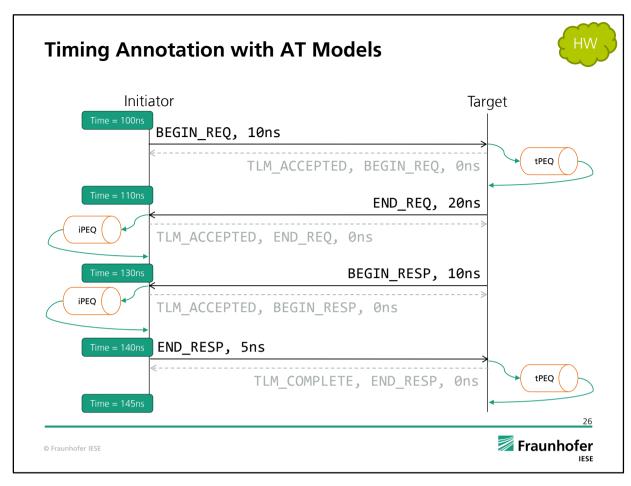


- AT models usually synchronize incoming calls with the simulation time
- Timing annotation is similar to b_transport:
 - Hey target, please pretend that you received this message 10ns in the future
- An AT component is usually not calling wait() statements for synchronizing with time, it rather posts the incoming transaction into a *Payload Event Queue* (PEQ). The PEQ internally synchronizes with the simulation time and calls a callback function in order to process the transaction when the time is reached.
- A component receives a transaction that should be processed in the future, so it puts the transaction into a PEQ in order to process it when time is reached

25



Your notes:		

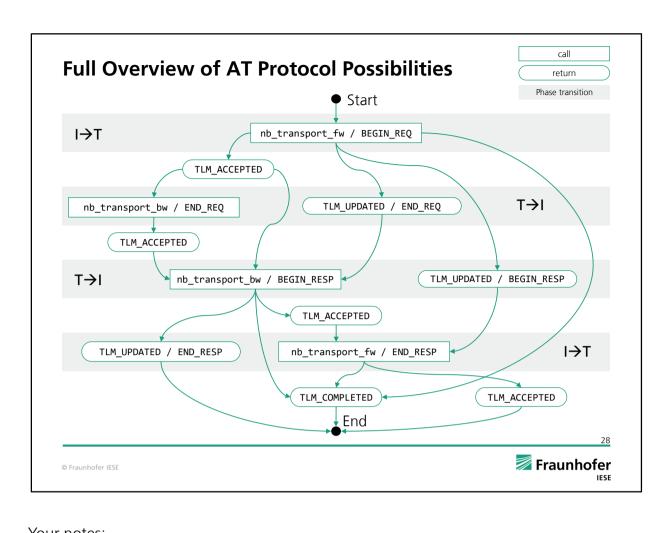


In AT models it is usual to synchronize incoming <code>nb_transport</code> calls with the simulation time. In the message sequence chart above, an initiator sends <code>BEGIN_REQ</code> with a timing annotation of 10 ns. The target returns <code>TLM_ACCEPTED</code> and is obliged to behave as if it had received the transaction at <code>sc_time_stamp() + 10ns</code>. An AT component is usually not calling wait statements for synchronizing with time, it rather posts the incoming transaction into a <code>Payload Event Queue</code> (PEQ). The payload event queue internally synchronizes with the simulation time (<code>wait</code>) and calls a callback function in order to process the transaction when the time is reached. In other words, a component receives a transaction that should be processed in the future, so it puts the transaction into a queue such that it is actually processed in the future.

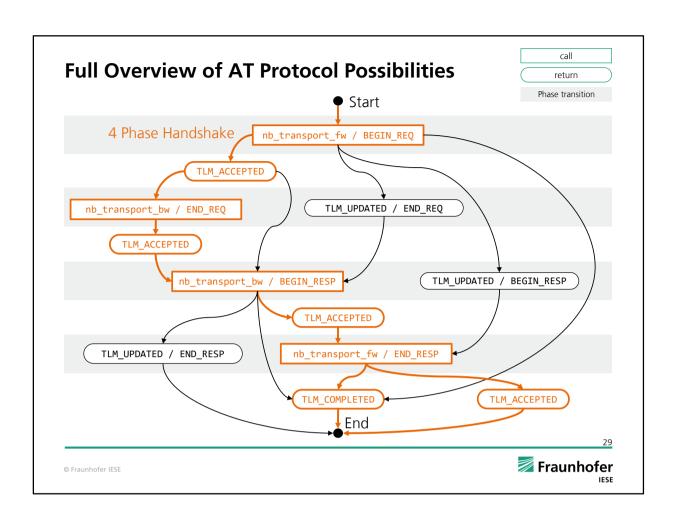
Your notes:			

					User L	ibraries			.	\'II.
Transaction Level Modeling (TLM)	Sockets & Generic	Blocking & Non-	Temporal Decoup- ling &	Phases	Payload Exten-	SystemC AMS	Electrical Linear Networks (ELN)		r Signal / (LSF)	Timed Data Flow (TDF)
saction	Payload	<u>Blocking</u>	DMI		sions	Syste	Linear D	AE solve	r	Scheduler
Trans							Sy	/nchroni:	zation laye	er
		Predefined Primitive Channels: Mutexes, FIFOs & Signals								
SystemC	Simulation		Methods & Threads		Channels & Interfaces			Data Types: gic, Integers,		
Kernel		Events, Sensitivity & Notifications			Fixe			dpoint & tingpoint		
					C	++				

Your notes:	



Your notes:		



Your notes:		

Initiator BEGIN_REQ, Ons TLM_ACCEPTED, BEGIN_REQ, Ons END_REQ, Ons BEGIN_RESP, Ons TLM_ACCEPTED, BEGIN_RESP, Ons TLM_ACCEPTED, BEGIN_RESP, Ons TLM_COMPLETE, END_RESP, Ons TLM_COMPLETE, END_RESP, Ons

Your notes:		

4 Phase Handshake: Initiator

```
HW
```

```
class Initiator: public sc_module, public tlm::tlm_bw_transport_if<> {
     public:
tlm::tlm_initiator_socket<> socket;
    MemoryManager mm;
int data[16];
tlm::tlm generic_payload* requestInProgress;
sc_event endNequest;
tlm_utils::peq_with_cb_and_phase<Initiator> peq;
    socket.hind(*this):
          SC_THREAD(process);
                                           Generate a sequence of
                                         random transactions
    protected:
void process() {
    tlm::tlm_generic_payload* trans;
    tlm::tlm_phase phase;
    sc_time delay;
                                                                Grab a new
                                                                transaction from
          for (int i = 0; i < 100; i++) {
    trans = mm.allocate();
    trans->acquire();
                                                                the memory
                                                                manager pool
               trans->set_command(...);
               trans->set_response_status(tlm::TLM_INCOMPLETE_RESPONSE);
               if (requestInProgress) {
   wait(endRequest);
                                                        BEGIN_REQ/END_REQ
                                                       exclusion rule
               }
requestInProgress = trans;
phase = tlm::BEGIN_REQ;
delay = ...;
               tlm:::tlm_sync_enum status;
status = socket->nb_transport_fw( *trans, phase, delay );
               wait(randomDelay());
```

```
peq.notify(trans, phase, delay);
return tlm::TLM_ACCEPTED;
                                 Queue the transaction into the peq
                                 until the annotated time has elapsed
                                                Payload event queue
caĺlback
       if (phase == tlm::END_REQ || ...)
          requestInProgress =
endRequest.notify();
                                                Wake-up suspended
                                               main process
       else if (phase == tlm::BEGIN_RESP)
                                               Do something with
          checkTransaction(trans);
          tlm::tlm_phase fw_phase = tlm::END_RESP; |
sc_time delay = sc_time(...);
                                               transaction
          socket->nb_transport_fw( trans, fw_phase, delay );
                             Allow MM to free the transaction object
       else if (phase == tlm::BEGIN_REQ || phase == tlm::END_RESP)
          SC_REPORT_FATAL(name(), "Illegal transaction phase received");
};
```



4 Phase Handshake: Target



```
class Target: public sc_module, public tlm::tlm_fw_transport_if<> {
   public:
   tlm::tim_target_socket<> socket;
   tlm::tlm_generic_payload* transactionInProgress;
   sc_event targetDone;
   bool responseInProgress;
   tlm::tlm_generic_payload* nextResponsePending;
   tlm::tlm_generic_payload* endRequestPending;
   tlm::tlm_generic_payload* endRequestPending;
   tlm:utils::peq_with_cb_and_phase<Target> peq;
                                                                                                                                 else if (phase == tlm::END_RESP) {
                                                                                                                                                                                            Flag must only be
                                                                                                                                                                                            cleared when
                                                                                                                                                                                            END RESP is sent
                                                                                                                                      responseInProgress = false;
                                                                                                                                      if (nextResponsePending) {
    sendResponse(*nextResponsePending);
      SC_CTOR(Target) : socket("socket"),
transactionInProgress(0),
responseInProgress(false),
nextResponsePending(0),
endRequestPending(0),
peq(this, &Target::peqCallback)
                                                                                                                                                                                                      Target itself is now
                                                                                                                                             nextResponsePending = 0;
                                                                                                                                                                                                      clear to issue the next
                                                                                                                                                                                                      BEGIN_RESP
                                                                                                                                      if (endRequestPending) {
    sendEndRequest(*endRequestPending);
    endRequestPending = 0;
             socket.bind(*this):
                                                                                                                                                                                                      unblock the initiator
                                                                                                                                                                                                      by issuing END_REQ
            SC_METHOD(executeTransactionProcess);
sensitive << targetDone; dont_initialize();</pre>
                                                                                                                                 else // tlm::END_REQ or tlm::BEGIN_RESP
                                                                                                                                      SC_REPORT_FATAL(name(), "Illegal transaction phase received");
      void sendEndRequest(tlm::tlm_generic_payload& trans)
            peq.notify( trans, phase, delay);
return tlm::TLM_ACCEPTED;
                                                                                                                                tlm::tlm_phase bw_phase;
sc_time delay;
                                                                                                                                 bw_phase = tlm::END_REQ;
delay = ...; // Accept delay
      tlm::tlm_sync_enum status;
status = socket->nb_transport_bw( trans, bw_phase, delay );
            sc time delay;
                                                                     Increment the
            if(phase == tlm::BEGIN_REQ) {
                                                                                                                                 delay = delay + randomDelay();
targetDone.notify( delay );
                                                                     transaction reference
             trans.acquire();
                                                                     count
                                                                                                                                 assert(transactionInProgress == 0);
transactionInProgress = &trans;
                  if (!transactionInProgress) {
    sendEndRequest(trans);
                                                                                                                                                                                           Queue internal event to mark
                  }
else {
   endRequestPending = &trans;
                                                                     Put back-pressure on
                                                                                                                                                                                           beginning of response
                                                                     initiator by deferring
                                                                     END REQ
```

Fraunhofer

4 Phase Handshake: Target



```
Method process that runs on
     void executeTransactionProcess()
{
                                                             targetDone event
          executeTransaction(*transactionInProgress);
                                                                        Target must honor
                                                                        BEGIN_RESP/END_RE
          if (responseInProgress)
                                                                        SP exclusion rule
               nextResponsePending = transactionInProgress;
               sendResponse(*transactionInProgress);
     }
     void sendResponse(tlm::tlm_generic_payload& trans)
                                                                        Function for sending
          tlm::tlm_sync_enum status;
tlm::tlm_phase bw_phase;
sc_time delay;
                                                                        the response
          responseInProgress = tr
          bw_phase = tlm::BEGIN_RESP;
delay = SC_ZERO_TIME;
status = socket->nb_transport_bw( trans, bw_phase, delay );
                                                                        Tell the memory
          trans.release();
                                                                        manager to free
                                                                        transaction
...
};
```

© Fraunhofer IESE

Fraunhofer

Base Protocol Rules [2]: Using the Return Paths



Initiator

Target

Time = 100ns

BEGIN_REQ, Ons

Time = 110ns

TIM UPDATED, END_REQ, 10ns

BEGIN_RESP, Ons

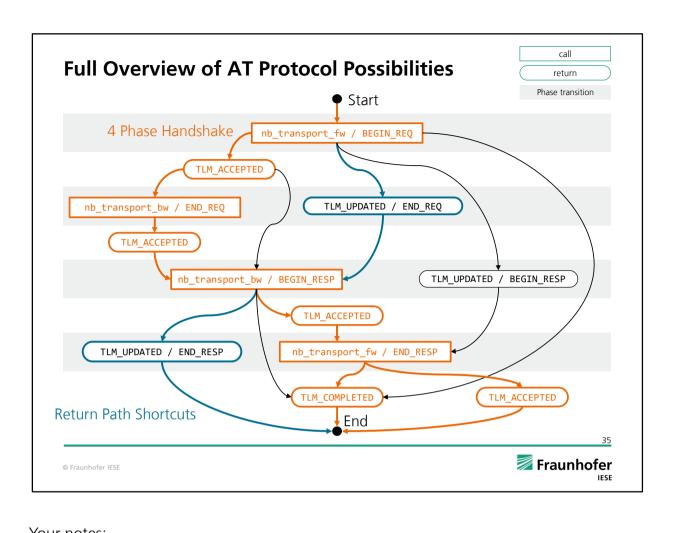
TLM UPDATED, END_RESP, 5ns

Time = 155ns

- The FW return path can be used as an alternative to the BW path
- The BW return path can be used as an alternative to the FW path
 - State must be set to TLM_UPDATED
- Timing must be increased



Your notes:		



Your notes:		

Base Protocol Rules [4]: Early Completion



Initiator

Target

Time = 100ns

BEGIN_REQ, Ons

TLM COMPLETED, d.c., 10ns

Time = 110ns

- The initiator sends BEGIN_REQ and the target immediately returns TLM_COMPLETED (useful for modelling simple I/O, where no ACK is required)
- The targets pre-empts any further communication by signaling a so called Early completion
- Initiator should immediacy check the response status of the generic payload object
- With TLM_COMPLETED the caller should always ignore the phase argument

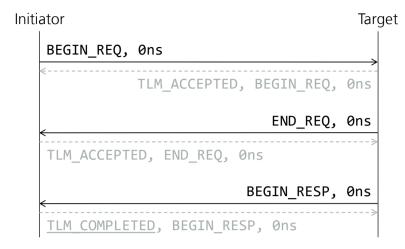
36



Your notes:			

Base Protocol Rules [4]: Early Completion



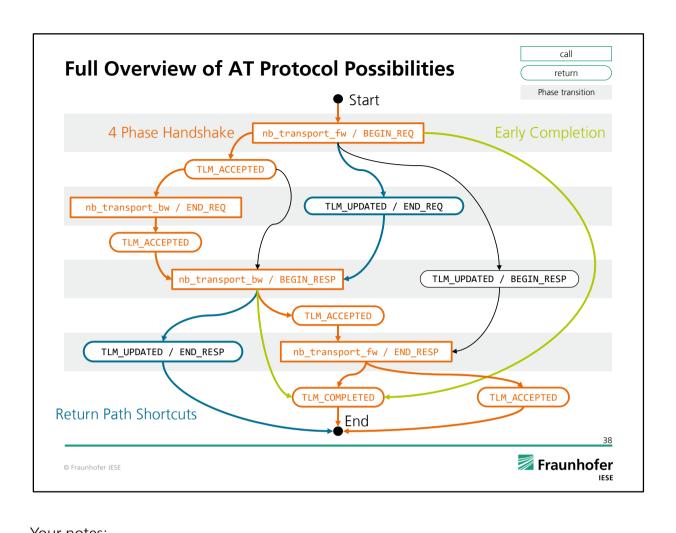


- The initiator can ealy complete the transaction by returning TLM_COMPLETED
- The initiator pre-empts therefore any further communication

37

Fraunhofer
IESE

ur notes:	



ur notes.

BEGIN_REQ, Ons TLM_ACCEPTED, BEGIN_RESP, Ons TLM_ACCEPTED, BEGIN_RESP, Ons TLM_ACCEPTED, BEGIN_RESP, Ons TLM_COMPLETED, END_RESP, Ons TLM_COMPLETED, END_RESP in the next BW call Fraunhofer IESE Your notes:

Your notes:			

Base Protocol Rules (7): Skip END_REQ (Shortcut) Initiator BEGIN_REQ, @ns TLM_UPDATED, BEGIN_RESP, @ ns END_REQ missed! END_RESP, @ns TLM_COMPLETED, END_RESP, @ns

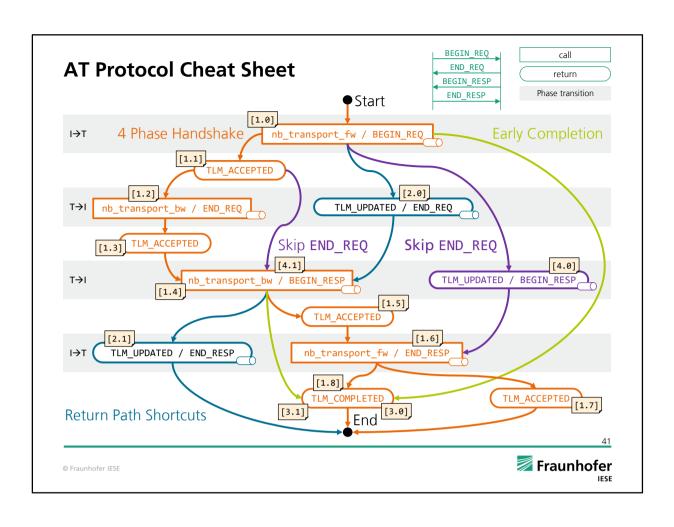
■ END_REQ is pre-empted by the target sending directly BEGIN_RESP via

TLM UPDATED over the return path

© Fraunhofer IESE

Your notes:		

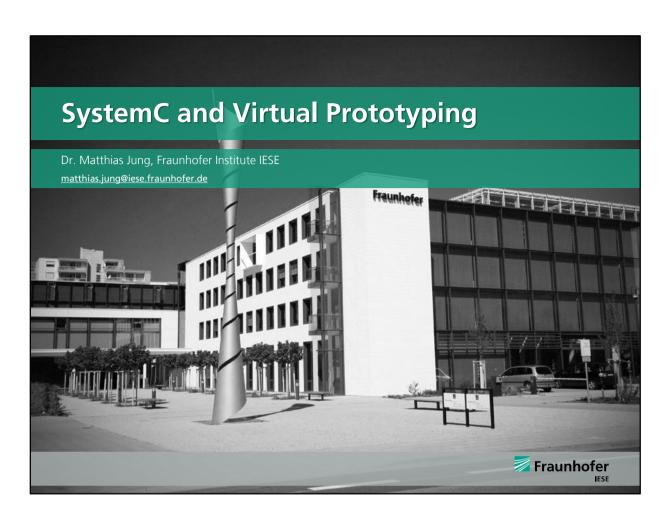
Fraunhofer



our notes:	

(s	Y S	ΓEΙ	ч с™		User L	ibraries				71001001110	
Transaction Level Modeling	Sockets & Generic	Blocking & Non-	Temporal Decoup- ling &	Phases	Payload Exten-	SystemC AMS	Electrical Linear Networks (ELN)		ear Signal ow (LSF)	Timed Data Flow (TDF)	8880
Transaction	Payload	Blocking	DMI		sions	Syste	Linear D Sy		ver onization laye	Scheduler	
	Predefined Primitive Channels: Mutexes, FIFOs & Signals									9	
SystemC	Simi	Simulation Kernel		Methods & Threads		Ch	annels & Interfac	es		a Types: , Integers,	•
	Kϵ			Events, Sensitivity & Notifications		Fix			dpoint & tingpoint		
C++											
All		0	- !		P. B. B. B. B.	A .		Š	#		

Your notes:			



Your notes:		



Your notes:	

					User L	braries			\\\\\\\\\\\\\\\\\\\\\\\\\\\\\\\\\\\\\\				
Transaction Level Modeling (TLM)	Sockets & Generic	Blocking & Non-	Temporal Decoup- ling &	Phases	Payload Exten-	SystemC AMS	Electrical Linear Networks (ELN)	Linear Signa Flow (LSF)	Timed Data Flow (TDF)				
saction	Payload	Blocking	DMI				sions		sions	Syste	Linear D.	AE solver	Scheduler
Tran		Synchronization layer					ayer						
		Predefined Primitive Channels: Mutexes, FIFOs & Signals											
SystemC	Simulation		Methods & Threads Simulation			Cł	annels & Interface		Data Types: gic, Integers,				
Ś	Kernel			Events, Sensitivity & Notifications			odules & Hierarch		ixedpoint & oatingpoint				
C++													

Your notes:	

Simple Sockets

- "Simple" because they are easy to use less to code ...
- Do not bind sockets to objects, instead register methods with each socket
- Dummy method implementations are provided:
 - get_direct_mem_ptr
 - transport_dbg
 - invalidate_direct_mem_ptr
 - nb_transport_bw
- Target need only register either b_transport or nb_transport_fw
- Automatic conversion between blocking and non-blocking

48



Your notes:	

Simple Sockets: Initiator Example #include "tlm_utils/simple_initiator_socket.h" Instantiate SC_MODULE(Initiator) { simple socket tlm_utils::simple_initiator_socket<Initiator> iSocket; iSocket.register_nb_transport_bw(this, &Initiator::nb_transport_bw); Register function SC_THREAD(process); } void process() { Implementation of Function } } Try code on github: https://github.com/TUK-SCVP/SCVP.artifacts/blob/master/tlm_simple_sockets/ }; Fraunhofer

Your notes:		

Simple Sockets: Target Example

© Fraunhofer IESE

```
#include "tlm_utils/simple_target_socket.h"
SC_MODULE(Target)
                                                                                   Instantiate simple
    tlm_utils::simple_target_socket<Target> tSocket;
                                                                                   target socket
    SC_CTOR(Target) : tSocket("tSocket"), ...
        tSocket.register_b_transport(this, &Target::b_transport);
                                                                                   Register functions
        tSocket.register_nb_transport_fw(this, &Target::nb_transport_fw);
    void b_transport(tlm::tlm_generic_payload& trans, sc_time& delay) {
                                                                                   Implementation of
                                                                                   Functions
    }
    tlm_sync_enum nb_transport_fw(tlm::tlm_generic_payload& ..., tlm::tlm_phase& ..., sc_time& ...) {
};
                                                                                        Fraunhofer
```

Tagged Simple Sockets: Interconnect Example

Initiator
A
Interconnect
A
Target
A
Target
A
Target
B

- simple_initiator_socket_tagged
- Register the same callback method with several sockets
- Distinguish origin of incoming transactions using socket id
- Interconnect is usually template class for number of target and initiator sockets



Your notes:		

Tagged Simple Sockets: Interconnect Example

```
template<unsigned int I, unsigned int T>
SC_MODULE(BUS)
    public:
    tlm_utils::simple_target_socket_tagged<BUS> tSocket[T];
    tlm_utils::simple_initiator_socket_tagged<BUS> iSocket[I];
         for(unsigned int i = 0; i < T; i++)</pre>
             for(unsigned int i = 0; i < I; i++)</pre>
             iSocket[i].register_nb_transport_bw(this,
    &BUS::nb_transport_bw, i);
    }
    private:
    std::map<tlm::tlm_generic_payload*, int> bwRoutingTable;
std::map<tlm::tlm_generic_payload*, int> fwRoutingTable;
 virtual void b_transport( int id,
                               tlm::tlm_generic_payload& trans,
                               sc_time& delay )
         sc_assert(id < T);
int outPort = routeFW(id, trans, false);</pre>
         iSocket[outPort]->b_transport(trans, delay);
```

```
virtual tlm::tlm_sync_enum nb_transport_fw( int id,
                   tlm::tlm_generic_payload& trans,
                   tlm::tlm phase& phase,
                  sc_time& delay )
         sc_assert(id < T);
int outPort = 0;</pre>
         if(phase == tlm::BEGIN_REQ) {
              trans.acquire();
outPort = routeFW(id, trans, true);
         else if(phase == tlm::END_RESP) {
   outPort = fwRoutingTable[&trans];
              trans.release();
         else {
              SC_REPORT_FATAL(name(), "ERROR!");
         return iSocket[outPort]->nb_transport_fw(
                                             trans, phase, delay);
    virtual tlm::tlm_sync_enum nb_transport_bw( int id,
                   tlm::tlm_generic_payload& trans,
                  tlm::tlm_phase& phase,
sc_time& delay )
         int inPort = bwRoutingTable[&trans];
         return tSocket[inPort]->nb_transport_bw(
                                              trans, phase, delay);
};
```

© Fraunhofer IESE



Simple Sockets: Target Example

© Fraunhofer IESE

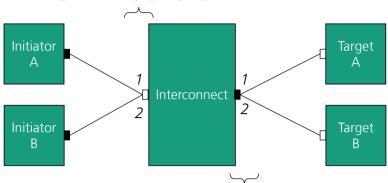
```
int routeFW(int inPort, tlm::tlm_generic_payload &trans, bool store)
    int outPort = 0;
    if(trans.get_address() < 512)</pre>
                                                                               Implementation of
                                                                               the memory map
        outPort = 0;
    else if(trans.get_address() >= 512 && trans.get_address() < 1024)</pre>
                                                                               Correct address,
                                                                               i.e. subtract offset
        trans.set_address(trans.get_address() - 512);
        outPort = 1;
                                                                               such that target
                                                                               gets always
                                                                               addresses starting
        trans.set_response_status( tlm::TLM_ADDRESS_ERROR_RESPONSE );
                                                                               at 0
    if(store) {
       bwRoutingTable[&trans] = inPort;
                                                                               Store from where
       fwRoutingTable[&trans] = outPort;
                                                                               transaction comes
                                                                               and where it goes
    return outPort;
                                                                                     Fraunhofer
```


Simple Sockets: Target Example

```
int sc_main (...)
                                                                                                   Number of
                                                                                                   components must
     Initiator * cpu1 = new Initiator("C1");
Initiator * cpu2 = new Initiator("C2");
                                                                                                   be known at
                                                                                                   compile time!
     Target * memory1 = new Target("M1");
Target * memory2 = new Target("M2");
     Interconnect<2,2> * bus = new BUS<2,2>("B1");
     cpu1->iSocket.bind(bus->tSocket[0]);
     cpu2->iSocket.bind(bus->tSocket[1]);
     bus->iSocket[0].bind(memory1->tSocket);
     bus->iSocket[1].bind(memory2->tSocket);
                                                                                Try code on github:
     sc_start();
                                                                              https://github.com/TUK-
                                                                  SCVP/SCVP.artifacts/blob/master/tlm_simple_sockets/
     return 0;
                                                                                                       Fraunhofer
© Fraunhofer IESE
```

Multipasstrough Sockets

Multi_passthrough_target_socket



multi_passtrough_initiator_socket

- Similar to Tagged sockets: also an id is used for identification
- The id is determined by the order of the binding to the multipasstrough socket
- Dynamic binding, number of comp. does not have to be known at compile time



Your notes:			

Multipasstrough Sockets: Interconnect Example

```
SC_MODULE(BUS)
     public:
     tlm_utils::multi_passthrough_target_socket<BUS> tSocket;
     tlm_utils::multi_passthrough_initiator_socket<BUS> iSocket;
    SC_CTOR(Interconnect) : tSocket("tSocket"), iSocket("iSocket")
         tSocket.register_b_transport(this, &BUS::b_transport);
         tSocket.register_nb_transport_fw(this, &BUS::nb_transport_fw); iSocket.register_nb_transport_bw(this, &BUS::nb_transport_bw);
    }
    std::map<tlm::tlm_generic_payload*, int> bwRoutingTable;
std::map<tlm::tlm_generic_payload*, int> fwRoutingTable;
     void b_transport( int id, tlm::tlm_generic_payload& trans, sc_time& delay ) {
    }
     tlm::tlm_sync_enum nb_transport_fw( int id, ... ) {
    }
    virtual tlm::tlm_sync_enum nb_transport_bw( int id, ...) {
    }
};
```

Your notes:

© Fraunhofer IESE

Fraunhofer

Multipasstrough Sockets: Interconnect Example

```
int sc_main (...)
     Initiator * cpu1 = new Initiator("C1");
     Initiator * cpu2 = new Initiator("C2");
     Target * memory1 = new Target("M1");
Target * memory2 = new Target("M2");
     Interconnect * bus = new BUS("B1");
                                                 The order of the
     cpu1->iSocket.bind(bus->tSocket);
                                                 binding
     cpu2->iSocket.bind(bus->tSocket);
                                                 determines the ids
                                                 for the function
     bus->iSocket.bind(memory1->tSocket);
                                                 calls!
     bus->iSocket.bind(memory2->tSocket);
     sc_start();
                                                           Try code on github:
     return 0;
                             https://github.com/TUK-SCVP/SCVP.artifacts/blob/master/tlm_multipasstrough_sockets
                                                                                                Fraunhofer
© Fraunhofer IESE
```

Your notes:			

	5)	/ S ⁻	ΓEΝ	d C™		User Li	braries	/=			31001001110	
ē,	Transaction Level Modeling (TLM)	Sockets & Generic	Blocking & <u>Non-</u>	Temporal Decoup- ling &	Phases	Payload Exten-	SystemC AMS	Electrical Linear Networks (ELN)		ar Signal ww (LSF)	Timed Data Flow (TDF)	00000
	ansaction	Payload	Blocking	DMI		sions	Syste	Linear D			Scheduler	
D ma		Synchronization layer Predefined Primitive Channels: Mutexes, FIFOs & Signals										
		Methods & Threads Simulation					Channels & Interfaces Data Types: Logic, Integers,					
vs vs	ίς.	Kε	ernel		Events, Sensitivity & Notifications					dpoint & tingpoint		
THE		WW. 27	Street April 1971			C-	++		and the second		III A	
	O.		0	. :		4	h . b .		Ď	d-h d-h		

Your notes:			

Modelling of Backpressure





- The exclusion rules enable flow control like back pressure:
 - E.g. if an input buffer of a target is full the target can defer the sending of END_REQ for the transaction that filled the buffer, until the buffer has available space again
 - An RTL ready signal can be modeled by deferring END_REQ
 - However, since it is non-blocking, the target can do something else, e.g. sending
- Also Response exclusion rule must be honored!

59



Your notes:		

Modelling of Backpressure



```
DECLARE_EXTENDED_PHASE(INTERNAL);
                                                                                                                                             Internal protocol phase, instead
      SC_MODULE(Target) {
                                                                                                                                             of using a sensitive process!
                   of using a set that it is the content of the content of using a set that it is the content of using a set that it is the content of using a set that it is the content of using a set to use the content of use the content
                    std::queue<tlm::tlm_generic_payload*> responseQueue;
                    void peqCallback(...) {
                                                                                                                                                                                                                    Check if
                                   if(phase == tlm::BEGIN_REQ) {
                                                                                                                                                                                                                    input buffer
                                                   trans.acquire();
if (numberOfTransactions < bufferSize) {
                                                                                                                                                                                                                    is full
                                                                 sendEndRequest(trans);
                                                  } else {
    endRequestPending = &trans;
                                                                                                                                                                                                                    Reduce
                                   } else if (phase == tlm::END_RESP) {
                                                                                                                                                                                                                    transaction
                                                                                                                                                                                                                    counter
                                                   numberOfTransactions--:
                                                   if (responseQueue.size() > 0) {
                                                              gp * next = responseQueue.front();
responseQueue.pop();
sendResponse(*next);
                                                                                                                                                                                                                    Send next
                                                                                                                                                                                                                    response
                                                   if (endRequestPending) {
    sendEndRequest(*endRequestPending);
    endRequestPending = 0;
                                                                                                                                                                                                                    Unblock
                                                                                                                                                                                                                  Initiator
                                    else if(phase == INTERNAL) {
                                                                                                                                                                                                                   Honor
                                                  executeTransaction(trans)
if (responseInProgress) {
                                                                                                                                                                                                                    response
                                                  responseQueue.push(&trans);
} else {
   sendResponse(trans);
                                                                                                                                                                                                                    exclusion
                                                                                                                                                                                                                   rule
© Fraunhofer IESE
```

Your notes.

Try code on github:

https://github.com/TUKSCVP/SCVP.artifacts/blob/master/tlm_at_backpressure

		100			User Li	braries		70.70		
Transaction Level Modeling (TLM)	Sockets & Generic	Blocking & Non-	Temporal Decoup- ling &	Phases	Payload Exten-	SystemC AMS	Electrical Linear Networks (ELN)		ar Signal w (LSF)	Timed Data Flow (TDF)
saction	Payload	<u>Blocking</u>	DMI		sions	Syste	Linear D.	AE solv	er	Scheduler
Trans							Sy	nchron/	ization laye	er
			Pred	defined Pri	mitive Cha	nnels:	Mutexes, FIFOs &	Signals		
SystemC	Simı	ulation	Me	thods & Th	nreads	Cł	annels & Interfac	es		a Types: , Integers,
S	Kε	ernel		nts, Sensit Notificatio		М	odules & Hierarch	ny		dpoint & tingpoint
Ì					C-	++				

Your notes:			

(s	Y	′ S ¯	ΓEΙ	M C™		User Li	braries	/=			31001001110	
Transaction Level Modeling	(TLM)	Sockets & Generic	Blocking & Non-	Temporal Decoup- ling &	Phases	Payload Exten-	SystemC AMS	Electrical Linear Networks (ELN)		ear Signal ow (LSF)	Timed Data Flow (TDF)	8880 JO.
ansaction I		Payload	Blocking	DMI		sions	Syste	Linear D			Scheduler	
E	+			Pred	defined Pri	mitive Cha	nnels:	S) Mutexes, FIFOs &		nization laye	er	
	oystellic.	Simı	ulation	T	thods & Th			nannels & Interfac		Dat	a Types:	•
	ń'	Ke	ernel		nts, Sensit Notificatic	•	М	odules & Hierarch	ny		dpoint & tingpoint	
THE STATE OF THE S	- Aug	MIOTALISM I				C-	++					
***			0	. !		A PARA	h .		b	d-h d-h		

		Your notes:

Payload Extensions

- Payload extensions are a flexible and powerful method to carry additional non-standard attributes (e.g. priority, ...)
- Extensions are attached to the normal generic payload
- Extensions can be:
 - Ignorable (ensures compatibility with base protocol)
 - Mandatory (a new protocol is needed!)
 - Private (only used by a single module e.g. Interconnect for storing routing information instead using maps)
 - Sticky (remain when transaction is returned to a pool)
 - Auto (freed when transaction is returned to a pool)

Generic payload object

Command Address Data Byte Enables Response Status

Extension A

Extension B

Extension C



64



Your notes:			
	 	 	· · · · · · · · · · · · · · · · · · ·

Payload Extensions Example

```
SC_MODULE (Interconnect)
{
    ...
    routingExtension* ext;
    ext = new routingExtension(
        inPort, outPort);
    trans.set_auto_extension(ext);
    ...
    routingExtension *ext = NULL;
    trans.get_extension(ext);
    outPort = ext->getOutputPortNumber();
    ...
};
```

Try code on github:
 https://github.com/TUKSCVP/SCVP.artifacts/blob/master/tlm_payload_extensions/



Your notes:			

(s)	ys-	ΓEΙ	ч с™		User L	ibraries				1001001110	
Transaction Level Modeling (TLM)	Sockets & Generic	Blocking & Non-	Temporal Decoup- ling &	Phases	Payload Exten-	SystemC AMS	Electrical Linear Networks (ELN)		ear Signal ow (LSF)	Timed Data Flow (TDF)	8880
Transaction (1	Payload	Blocking	DMI		sions	Syste	Linear D Sy		lver onization laye	Scheduler	
			Pred	defined Pri	mitive Cha	nnels: I	Mutexes, FIFOs &	Signa	ls		19
SystemC	Simu	ulation	Me	thods & Th	nreads	Ch	annels & Interfac	es		a Types: , Integers,	•
	Кє	ernel		nts, Sensit Notificatio		M	odules & Hierarch	ıy		dpoint & tingpoint	
THE					С	++					
411		0	. !		P. B. B. B.	h :		Ď.	##		

Your notes:		